



**Due date: Monday September 18th,  
11:59pm**

Page 1 of 3

The main program `mp2.py` is provided to help you test your program. The autograder does not use it. So feel free to modify the default values of tunable parameters (near the end of the file). Do not modify `reader.py`.

To run the main program, type **`python3 mp2.py`** in your terminal. This should load the provided datasets and run the `bigramBayes` function on them. Sadly, it's not doing the required training and just returns the label -1 for all the reviews. Your job is to write real code for `bigramBayes`, returning a list of 0's (Negative) and 1's (Positive).

The main program `mp2.py` accepts values for a number of tunable parameters. To see the details, type **`python3 mp2.py -h`** in your terminal. Note that you can and should change the parameters as necessary to achieve good performance.

## Submitting to Gradescope

Submit this assignment by uploading `bigram_naive_bayes.py` to Gradescope. You can upload other files with it, but only `bigram_naive_bayes.py` will be retained by the autograder. We strongly encourage you to submit to Gradescope early and often as you will be able to see your final score there.

## Policies

You are expected to be familiar with the general policies on the course syllabus (e.g. academic integrity) and on the top-level MP page (e.g. code style). In particular, notice that this is an individual assignment.

## Dataset

The dataset in your template package consists of 10000 positive and 3000 negative movie reviews. It is a subset of the [Stanford Movie Review Dataset](#), which was originally introduced by [this paper](#). We have split this data set for you into 5000 development examples and 8000 training examples. The autograder also has a hidden set of test examples, generally similar to the development dataset.

## Bigram Mixture Model

You will implement the function `bigramBayes` that computes the mixture of unigram and bigram bag of words models. Each bigram  $b_i$  is a sequence of two consecutive words from a training or test review. Your bigram code should be very similar to your unigram code, except that you're looking at pairs of words rather than single words. So the probabilities for bigram and unigram models look like this:

$$P(\text{Type} = \text{Positive} | \text{Words}) = \frac{P(\text{Type} = \text{Positive})}{P(\text{Words})} \prod_{\text{All word pairs}} P(\text{Word Pair} | \text{Type} = \text{Positive})$$

$$P(\text{Type} = \text{Negative} | \text{Words}) = \frac{P(\text{Type} = \text{Negative})}{P(\text{Words})} \prod_{\text{All word pairs}} P(\text{Word Pair} | \text{Type} = \text{Negative})$$

Then you combine the bigram model and the unigram model into a mixture model defined with parameter  $\lambda$ :

$$(1 - \lambda) \log \left[ P(Y) \prod_{i=1}^n P(w_i|Y) \right] + \lambda \log \left[ P(Y) \prod_{i=1}^m P(b_i|Y) \right]$$

The input to `bigramBayes` includes two Laplace smoothing parameters, one for the unigram model and one for the bigram model.

The parameter  $\lambda$  controls how much emphasis to give to the unigram model and how much to the bigram model. Choose the value of  $\lambda$  that gives the highest classification accuracy and set this to be the default value of the parameter in `bigramBayes`.

As with MP1, some of our autograder tests will use your default values for the tunable parameters inputs to `bigramBayes`. However, some of our tests will reset these parameter values so that we can test specific aspects of your code. In particular, your definition of `bigramBayes` should set the default value of `pos_prior` to match the development dataset and we will adjust this value for our hidden datasets.

You can continue to experiment with stemming, transforming to lowercase, and/or removing stop words. If you turn on these features from the command line or in `load_data`, they will be used by `bigramBayes`. You can also use them by editing `bigramBayes`.

## Making the details work

Consider Python's Counter data structure.

**Use the log of the probabilities to prevent underflow/precision issues.** Apply log to both sides of the equation and convert multiplication to addition. Be aware that the standard python math functions are faster than the corresponding numpy functions, when applied to individual numbers.

Zero values in the naive Bayes equations will prevent the classification from working right. Therefore, you must smooth your calculated probabilities so that they are never zero. In order to accomplish this task, use Laplace smoothing. See the lecture notes for details. The Laplace smoothing parameter  $\alpha$  is passed as an argument to `bigramBayes` and you can adjust its value using the command-line arguments to `mp2.py`.

Tune the values of the Laplace smoothing constant using the command-line arguments to `mp2.py`. When you are happy with the result on the development set, edit the default values for these parameters in the definition of the function `bigramBayes`. Some of our tests will use your default settings and some tests will pass in new values.

You can experiment with other methods that might (or might not) improve performance. The command line options will let you transform the input words by converting them all to lowercase and/or running them through the Porter Stemmer. If you wish to turn either of these on for your autograder tests, edit the default values in the function `load_data`.

You could also try removing stop words from the reviews before you process them. You can add this to `load_data` or to the start of your `bigramBayes` function. You will need to find a suitable list of stop words and write a short python function to modify the input data.

No guarantees about what changes will make the accuracy better or worse. You need to figure that out by experimenting.